



Reduced-model scheduling policy comparison

Profile: Measured NVIDIA L4 profile for Qwen/Qwen2.5-7B-Instruct served by pinned vLLM


exact dp



broad fqi



narrow fqi



decode jobs d

prefill (P)

decode (D)

idle (I)

logged visits to displayed action