

Requests, work, and the GPU state

Profile: Measured NVIDIA L4 profile for Qwen/Qwen2.5-7B-Instruct served by pinned vLLM

Maximum clock

unfinished requests

realized clock (MHz)

time (s)

