



Reduced-model prefill and decode scheduling

Profile: Measured NVIDIA L4 profile for Qwen/Qwen2.5-7B-Instruct served by pinned vLLM


oldest age bin 0



oldest age bin 2



oldest age bin 4



prefill (P) decode (D) idle (I)